

Q: What is the difference between inner and outer join? Explain with example.

Inner join

Inner join is the most common type of Join which is used to combine the rows from two tables and create a result set containing only such records that are present in both the tables based on the joining condition (predicate).

Inner join returns rows when there is at least one match in both tables

If none of the record matches between two tables, then INNER JOIN will return a NULL set. Below is an example of INNER JOIN and the resulting set.

Outer Join

Outer Join can be full outer or single outer

Outer Join, on the other hand, will return matching rows from both tables as well as any unmatched

rows from one or both the tables (based on whether it is single outer or full outer join respectively).

Notice in our record set that there is no employee in the department 5 (Logistics). Because of this if we perform inner join, then Department 5 does not appear in the above result. However in the below query we perform an outer join (dept left outer join emp), and we can see this department.

INNER JOIN: gets all records from one table that have some related entry in a second table

LEFT JOIN: gets all records from the LEFT linked table but if you have selected some columns from the RIGHT table, if there is no related records, these columns will contain NULL

RIGHT JOIN: is like the above but gets all records in the RIGHT table

FULL JOIN: gets all records from both tables and puts NULL in the

columns where related records do not exist in the opposite table

Q: What is the difference between JOIN and UNION?

SQL JOIN allows us to “lookup” records on other table based on the given conditions between two tables. For example, if we have the department ID of each employee, then we can use this department ID of the employee table to join with the department ID of department table to lookup department names.

UNION operation allows us to add 2 similar data sets to create resulting data set that contains all the data from the source data sets. Union does not require any condition for joining. For example, if you have 2 employee tables with same structure, you can UNION them to create one result set that will contain all the employees from both of the tables.

SELECT \* FROM EMP1

UNION

SELECT \* FROM EMP2;

What is the difference between UNION and UNION ALL?

UNION removes duplicate records (where all columns in the results are the same), UNION ALL does not.

There is a performance hit when using UNION vs UNION ALL, since the database server must do additional work to remove the duplicate rows, but usually you do not want the duplicates (especially when developing reports).

UNION Example:

SELECT 'foo' AS bar UNION SELECT 'foo' AS bar

Result:

+-----+

| bar |

+-----+

| foo |

+-----+

1 row in set (0.00 sec)

UNION ALL example:

SELECT 'foo' AS bar UNION ALL  
SELECT 'foo' AS bar

Result:

+-----+

| bar |

+-----+

| foo |

| foo |

+-----+

2 rows in set (0.00 sec)

Q: What is the difference between WHERE clause and HAVING clause?

WHERE clause introduces a condition on individual rows;  
HAVING clause introduces a

condition on aggregations, i.e. results of selection where a single result, such as count, average, min, max, or sum, has been produced from multiple rows. Your query calls for a second kind of condition (i.e. a condition on an aggregation) hence HAVING works correctly.

As a rule of thumb, use WHERE before GROUP BY and HAVING after GROUP BY. It is a rather primitive rule, but it is useful in more than 90% of the cases.

While you're at it, you may want to re-write your query using ANSI version of the join:

SELECT L.LectID,Fname,Lname

FROM Lecturers L

JOIN Lecturers\_Specialization S ON  
L.LectID=S.LectID

GROUP BY L.LectID,Fname,Lname

HAVING COUNT(S.Expertise)>=ALL

(SELECT COUNT(Expertise) FROM Lecturers\_Specialization GROUP BY LectID)

This would eliminate WHERE that was used as a theta join condition.

Q: What is the difference among UNION, MINUS and INTERSECT?

UNION combines the results from 2 tables and eliminates duplicate records from the result set.

MINUS operator when used between 2 tables, gives us all the rows from the first table except the rows which are present in the second table.

INTERSECT operator returns us only the matching or common rows between 2 result sets.

To understand these operators, let's see some examples. We will use two different queries to extract data from our emp table and then we will perform UNION, MINUS and

INTERSECT operations on these two sets of data.

UNION

SELECT \* FROM EMPLOYEE WHERE ID = 5

UNION

SELECT \* FROM EMPLOYEE WHERE ID = 6

ID	MGR_ID	DEPT_ID	NAME	SAL	DOJ
5	2	2.0	Anno	80.0	01-Feb-2012
6	2	2.0	Darl	80.0	11-Feb-2012

MINUS

SELECT \* FROM EMPLOYEE

MINUS

SELECT \* FROM EMPLOYEE WHERE ID > 2

ID	MGR_ID	DEPT_ID	NAME	SAL	DOJ
----	--------	---------	------	-----	-----

1	2	Hash	100.0	01-Jan-2012
---	---	------	-------	-------------

2	1	2	Robo	100.0	01-Jan-2012
---	---	---	------	-------	-------------

INTERSECT

SELECT \* FROM EMPLOYEE WHERE ID IN (2, 3, 5)

INTERSECT

SELECT \* FROM EMPLOYEE WHERE ID IN (1, 2, 4, 5)

ID	MGR_ID	DEPT_ID	NAME	SAL	DOJ
5	2	2	Anno	80.0	01-Feb-2012
2	1	2	Robo	100.0	01-Jan-2012

Q: How to generate row number in SQL Without ROWNUM

Generating a row number – that is a running sequence of numbers for each row is not easy using plain SQL. In fact, the method I am

going to show below is not very generic either. This method only works if there is at least one unique column in the table. This method will also work if there is no single unique column, but collection of columns that is unique. Anyway, here is the query:

```
SELECT name, sal, (SELECT  
COUNT(*) FROM EMPLOYEE i  
WHERE o.name >= i.name)  
row_num
```

```
FROM EMPLOYEE o
```

```
order by row_num
```

Q: How to select first 5 records from a table?

This question, often asked in many interviews, does not make any sense to me. The problem here is how do you define which record is first and which is second. Which record is retrieved first from the database is not deterministic. It depends on many uncontrollable factors such as how database works

at that moment of execution etc. So the question should really be – “how to select any 5 records from the table?” But whatever it is, here is the solution:

In Oracle,

```
SELECT *
```

```
FROM EMP
```

```
WHERE ROWNUM <= 5;
```

In SQL Server,

```
SELECT TOP 5 * FROM EMP;
```

Generic solution,

I believe a generic solution can be devised for this problem if and only if there exists at least one distinct column in the table. For example, in our EMP table ID is distinct. We can use that distinct column in the below way to come up with a generic solution of this question that does not require database specific functions such as ROWNUM, TOP etc.

```
SELECT name
```

```
FROM EMPLOYEE o
```

```
WHERE (SELECT count(*) FROM  
EMPLOYEE i WHERE i.name <  
o.name) < 5
```

What is the difference between ROWNUM pseudo column and ROW\_NUMBER() function?

ROWNUM is a pseudo column present in Oracle database returned result set prior to ORDER BY being evaluated. So ORDER BY ROWNUM does not work.

ROW\_NUMBER() is an analytical function which is used in conjunction to OVER() clause wherein we can specify ORDER BY and also PARTITION BY columns.

Suppose if you want to generate the row numbers in the order of ascending employee salaries for example, ROWNUM will not work. But you may use ROW\_NUMBER() OVER() like shown below:

```
SELECT name, sal, row_number()
over(order by sal desc)
rownum_by_sal
```

```
FROM EMPLOYEE o
```

What are the differences among  
ROWNUM, RANK and  
DENSE\_RANK?

ROW\_NUMBER assigns contiguous,  
unique numbers from 1.. N to a  
result set.

RANK does not assign unique  
numbers—nor does it assign  
contiguous numbers. If two records  
tie for second place, no record will  
be assigned the 3rd rank as no one  
came in third, according to RANK.  
See below:

```
SELECT name, sal, rank()
over(order by sal desc) rank_by_sal
FROM EMPLOYEE o
```

name	Sal	RANK_BY_SAL
Hash	100	1

DENSE\_RANK, like RANK, does not  
assign unique numbers, but it does  
assign contiguous numbers. Even  
though two records tied for second  
place, there is a third-place record.  
See below:

```
SELECT name, sal, dense_rank()
over(order by sal desc)
dense_rank_by_sal
```

```
FROM EMPLOYEE o
```

name	Sal	DENSE_RANK_BY_SAL
Hash	100	1
Robo	100	1
Anno	80	2
Darl	80	2
Tomiti	70	3
Pete	70	3
Bhuti	60	4
Meme	60	4
Inno	50	5
Privy	50	5

Q: Clustered Indexes

Clustered indexes are indexes that  
uniquely identify the rows in an SQL  
table.

Every table can have exactly one  
clustered index.

You can create a clustered index  
that covers more than one column.  
For example: create Index  
index\_name(col1, col2, col.....).

By default, a column with a primary key already has a clustered index.

Q: Non-clustered Indexes

Non-clustered indexes are like simple indexes. They are just used for fast retrieval of data. Not sure to have unique data.

Q: What is the difference between a HAVING CLAUSE and a WHERE CLAUSE?

They specify a search condition for a group or an aggregate. But the difference is that HAVING can be used only with the SELECT statement. HAVING is typically used in a GROUP BY clause. When GROUP BY is not used, HAVING behaves like a WHERE clause. Having Clause is basically used only with the GROUP BY function in a query whereas WHERE Clause is applied to each row before they are part of the GROUP BY function in a query.

Q: What is Stored Procedure?

A stored procedure is a named group of SQL statements that have been previously created and stored in the server database. Stored procedures accept input parameters so that a single procedure can be used over the network by several clients using different input data. And when the procedure is modified, all clients automatically get the new version. Stored procedures reduce network traffic and improve performance. Stored procedures can be used to help ensure the integrity of the database.

Q: What are DMVs?

Dynamic Management Views (DMVs), are functions that give you information on the state of the server. DMVs, for the most part, are used to monitor the health of a server. They really just give you a snapshot of what's going on inside the server. They let you monitor the health of a server instance, troubleshoot major problems and

tune the server to increase performance.

Q: Define a temp table

In a nutshell, a temp table is a temporary storage structure. What does that mean? Basically, you can use a temp table to store data temporarily so you can manipulate and change it before it reaches its destination format.

Q: What's the difference between a local temp table and a global temp table?

Local tables are accessible to a current user connected to the server. These tables disappear once the user has disconnected from the server. Global temp tables, on the other hand, are available to all users regardless of the connection. These tables stay active until all the global connections are closed.

Q: How do you use transactions?

In general, there are three types of transactions that you can use in the SQL Server environment: BEGIN TRANSACTION, ROLL BACK TRANSACTION and COMMIT TRANSACTION. The gist behind deploying transactions is that they allow you to group multiple SQL commands into a single unit. From there, each transaction begins with a certain task, and ends when all the tasks within the transaction are complete. BEGIN TRANSACTION gets the ball rolling. ROLLBACK TRANSACTION functions a lot like an “undo” command, and COMMIT TRANSACTION completes all of the tasks within that transaction.

Q: What’s the difference between a clustered and a non-clustered index?

A clustered index directly affects the way tabled data is stored on a specific disk. This means that when a clustered index is used, data is stored in sequential rows based on the index column value. This is why

a table can only contain a single clustered index. Non-clustered indexes directly affect the way physical data is stored and managed within SQL Server.

Q: What are DBCC commands?

In very basic terms the Database Consistency Checker (DBCC) is used to aid in server maintenance. DBCC commands, many of which are completely undocumented, provide a set of commands that let you perform routine maintenance, status and validation checks. The most common DBCC commands are: DBCC CHECKALLOC (Lets you check disk allocation); DBCC OPENTRAN (Lets you check any open transactions); and DBCC HELP (shows a list of available DBCC commands to aid your server maintenance processes).

Q: Describe the difference between truncate and delete

The difference between these two processes is fairly simple. Truncate means to simply empty out a table. On the other hand, the delete command lets you delete entire rows from within a table, but not all of the data within that table.

Q: What is a view?

A view is simply a virtual table that is made up of elements of multiple physical or “real” tables. Views are most commonly used to join multiple tables together, or control access to any tables existing in background server processes.

Q: What is a Query Execution Plan?

SQL Server has several built-in tools that optimize how queries are executed within their databases. A query execution plan is exactly what it sounds like – a snapshot of how the optimizing tools will execute and deploy specific queries within the database. This service helps you troubleshoot problems

with jobs that don't necessarily execute perfectly.

Q: What is the default port number for SQL Server?

While this is kind of a softball question – if you know anything about SQL Server you should at least know the basic configuration options – it's an important one to nail in the interview. Basically, when SQL Server is enabled the server instance listens to the TCP port 1433.

Q: What is Cursor?

Cursor is a database object used by applications to manipulate data in a set on a row-by-row basis, instead of the typical SQL commands that operate on all the rows in the set at one time.

In order to work with a cursor we need to perform some steps in the following order:

Declare cursor

Open cursor

Fetch row from the cursor

Process fetched row

Close cursor

Deallocate cursor

What are the different index configurations a table can have?

A table can have one of the following index configurations:

No indexes

A clustered index

A clustered index and many nonclustered indexes

A nonclustered index

Many nonclustered indexes

What are different types of Collation Sensitivity?

Case sensitivity - A and a, B and b, etc.

Accent sensitivity

Kana Sensitivity - When Japanese kana characters Hiragana and Katakana are treated differently, it is called Kana sensitive.

Width sensitivity - A single-byte character (half-width) and the same character represented as a double-byte character (full-width) are treated differently than it is width sensitive.

Q: What are the properties of the Relational tables?

Relational tables have six properties:

1. Values are atomic.
2. Column values are of the same kind.
3. Each row is unique.
4. The sequence of columns is insignificant.



5. The sequence of rows is insignificant.

6. Each column must have a unique name.

You want to implement the following relationships while designing tables. How would you do it?

a.) One-to-one

b.) One-to-many

c.) Many-to-many

a.) One-to-One relationship - can be implemented as a single table and rarely as two tables with primary and foreign key relationships.

b.) One-to-Many relationships - by splitting the data into two tables with primary key and foreign key relationships.

c.) Many-to-Many - by using a junction table with the keys from both the tables forming the

composite primary key of the junction table.

Q: What is a Trigger

A trigger is a special kind of a store procedure that executes in response to certain action on the table like insertion, deletion or updation of data. It is a database object which is bound to a table and is executed automatically. You can't explicitly invoke triggers. The only way to do this is by performing the required action on the table that they are assigned to.

Types Of Triggers

After Triggers (For Triggers)

Instead Of Triggers

After Triggers

These triggers run after an insert, update or delete on a table. They are not supported for views.

AFTER INSERT Trigger.

AFTER UPDATE Trigger.

AFTER DELETE Trigger.

Instead Of Triggers

These can be used as an interceptor for anything that anyone tried to do on our table or view. If you define an Instead Of trigger on a table for the Delete operation, they try to delete rows, and they will not actually get deleted (unless you issue another delete instruction from within the trigger)

INSTEAD OF INSERT Trigger.

INSTEAD OF UPDATE Trigger.

INSTEAD OF DELETE Trigger.

Q: Difference between Stored Procedure and Function in SQL Server

Basic Difference

Function must return a value but in Stored Procedure it is optional(

Procedure can return zero or n values).

Functions can have only input parameters for it whereas Procedures can have input/output parameters .

Functions can be called from Procedure whereas Procedures cannot be called from Function.

#### Advance Difference

Procedure allows SELECT as well as DML(INSERT/UPDATE/DELETE) statement in it whereas Function allows only SELECT statement in it.

Procedures can not be utilized in a SELECT statement whereas Function can be embedded in a SELECT statement.

Stored Procedures cannot be used in the SQL statements anywhere in the WHERE/HAVING/SELECT section whereas Function can be.

The most important feature of stored procedures over function is

to retention and reuse the execution plan while in case of function it will be compiled every time.

Functions that return tables can be treated as another rowset. This can be used in JOINS with other tables.

Inline Function can be thought of as views that take parameters and can be used in JOINS and other Rowset operations.

Exception can be handled by try-catch block in a Procedure whereas try-catch block cannot be used in a Function.

We can go for Transaction Management in Procedure whereas we can't go in Function.

Q: Differentiate between a Local and a Global temporary table?

- A local temporary table exists only for the duration of a connection or, if defined inside a compound

statement, for the duration of the compound statement.

- Global temporary tables (created with a double "##") are visible to all sessions.

- Global temporary tables are dropped when the session that created it ends, and all other sessions have stopped referencing it.

Q: What are the properties and different Types of Sub-Queries?

#### Properties of Sub-Query

A sub-query must be enclosed in the parenthesis.

A sub-query must be put in the right hand of the comparison operator, and

A sub-query cannot contain an ORDER-BY clause.

A query can contain more than one sub-query.

## Types of Sub-Query

Single-row sub-query, where the sub-query returns only one row.

Multiple-row sub-query, where the sub-query returns multiple rows, and

Multiple column sub-query, where the sub-query returns multiple columns

Q: What is SQL Profiler?

SQL Profiler is a graphical tool that allows system administrators to monitor events in an instance of Microsoft SQL Server. You can capture and save data about each event to a file or SQL Server table to analyze later. For example, you can monitor a production environment to see which stored procedures are hampering performances by executing too slowly.

Use SQL Profiler to monitor only the events in which you are interested.

If traces are becoming too large, you can filter them based on the information you want, so that only a subset of the event data is collected. Monitoring too many events adds overhead to the server and the monitoring process and can cause the trace file or trace table to grow very large, especially when the monitoring process takes place over a long period of time.

Name 3 ways to get an accurate count of the number of records in a table?

```
SELECT * FROM table1
```

```
SELECT COUNT(*) FROM table1
```

```
SELECT rows FROM sysindexes  
WHERE id = OBJECT_ID(table1)  
AND indid < 2
```

Q: What is the STUFF function and how does it differ from the REPLACE function?

STUFF function is used to overwrite existing characters. Using this

syntax, STUFF (string\_expression, start, length, replacement\_characters), string\_expression is the string that will have characters substituted, start is the starting position, length is the number of characters in the string that are substituted, and replacement\_characters are the new characters interjected into the string. REPLACE function to replace existing characters of all occurrences. Using the syntax REPLACE (string\_expression, search\_string, replacement\_string), where every incidence of search\_string found in the string\_expression will be replaced with replacement\_string.

Q: What is CHECK Constraint?

A CHECK constraint is used to limit the values that can be placed in a column. The check constraints are used to enforce domain integrity.

How to get @@ERROR and @@ROWCOUNT at the same time?

If @@Rowcount is checked after Error checking statement then it will have 0 as the value of @@Recordcount as it would have been reset. And if @@Recordcount is checked before the error-checking statement then @@Error would get reset. To get @@error and @@rowcount at the same time do both in same statement and store them in local variable.

```
SELECT @RC = @@ROWCOUNT,  
@ER = @@ERROR
```

Q: What is maximum size row

In SQL 2000, the row limit is 8K bytes, which is the same size as a page in memory.

In 2005, the page size is the same (8K), but the database uses pointers on the row in the page to point to other pages that contain larger fields. This allows 2005 to overcome the 8K row size limitation.

Q: Query to get nth(3rd) Highest Salary

Select TOP 1 Salary as '3rd Highest Salary'

from (SELECT DISTINCT TOP 3  
Salary from Employee ORDER BY  
Salary DESC)

a ORDER BY Salary ASC

Q: Query to get nth(3rd) Lowest Salary

Select TOP 1 Salary as '3rd Lowest Salary'

from (SELECT DISTINCT TOP 3  
Salary from Employee ORDER BY  
Salary ASC)

a ORDER BY Salary DESC